


INF5062 – GPU & CUDA

Håkon Kvale Stensland
Simula Research Laboratory

PC Graphics Timeline

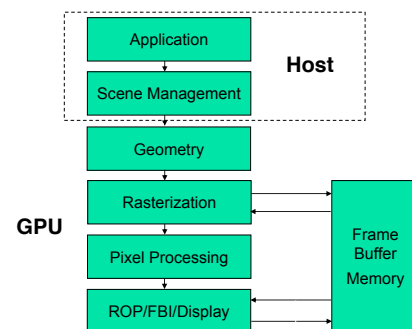
- Challenges:
 - Render infinitely complex scenes
 - And extremely high resolution
 - In 1/60th of one second (60 frames per second)
- Graphics hardware has evolved from a simple hardwired pipeline to a highly programmable multiword processor



(Some) 3D Buzzwords!!


- GPU: Graphics Processing Unit: A graphics chip with integrated programmable geometry and pixel processing
- Fill Rate: How fast the GPU can generate pixels, often a strong predictor for application frame rate
- Shader: A set of software instructions, which is used by the graphic resources primarily to perform rendering effects.
- API: Application Programming Interface – The standardized layer of software that allows applications (like games) to talk to other software or hardware to get services or functionality from them – such as allowing a game to talk to a graphics processor
- DirectX: Microsoft's API for media functionality
- Direct3D: Portion of the DirectX API suite that handles the interface to graphics processors
- OpenGL: An open standard API for graphics functionality. Available across platforms. Popular with workstation applications

Basic 3D Graphics Pipeline



Graphics in the PC Architecture

- FSB connection between processor and Northbridge (P45)
 - Memory Control Hub
- Northbridge handles PCI Express 2.0 to GPU and DRAM.
 - PCIe 2 x16 bandwidth at 16 GB/s (8 GB in each direction)
- Southbridge (ICH10) handles all other peripherals



High-end Hardware



- nVidia GeForce GTX 280
- Based on the latest generation GPU, codenamed GT200
- 1400 million transistors
- 240 Processing cores (SP) at 1296MHz
- 1024 MB Memory with 141.7GB/sec of bandwidth.
- 933 GFLOPS of computing power

Lab Hardware

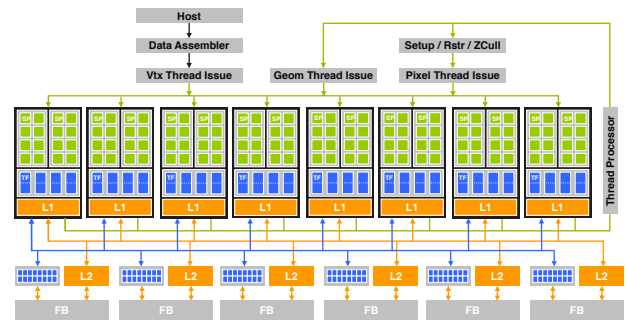


- nVidia GeForce 8600GT
- Based on the G84 chip
 - 289 million transistors
 - 32 Processing cores (SP) at 1190MHz
 - 512/256 MB Memory with 22.4GB/sec bandwidth

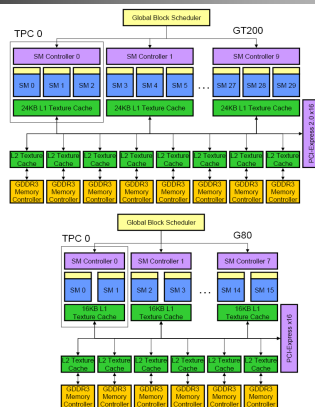


- nVidia GeForce 8800GT
- Based on the G92 chip
 - 754 million transistors
 - 112 Processing cores (SP) at 1500MHz
 - 256 MB Memory with 57.6GB/sec bandwidth


GeForce G80 Architecture



nVIDIA G80 vs. GT92/GT200 Architecture

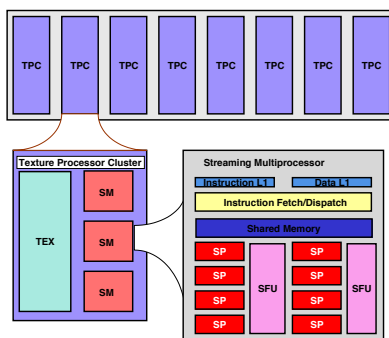


Compared with a multicore RISC-CPU




TPC... SM... SP... Some more details...

- TPC
 - Texture Processing Cluster
- SM
 - Streaming Multiprocessor
 - In CUDA: Multiprocessor, and fundamental unit for a thread block
- TEX
 - Texture Unit
- SP
 - Stream Processor
 - Scalar ALU for single CUDA thread
- SFU
 - Super Function Unit




SP: The basic processing block

- The nVIDIA Approach:
 - A Stream Processor works on a single operation
- AMD GPU's work on up to five operations, and Intel's Larrabee will work on up to 16
- Now, let's take a step back for a closer look!



Streaming Multiprocessor (SM)


- Streaming Multiprocessor (SM)
 - 8 Streaming Processors (SP)
 - 2 Super Function Units (SFU)
- Multi-threaded instruction dispatch
 - 1 to 768 threads active
 - Try to Cover latency of texture/memory loads
- Local register file (RF)
- 16 KB shared memory
- DRAM texture and memory access



Foils adapted from nVIDIA


SM Register File

- Register File (RF)
 - 32 KB
 - Provides 4 operands/clock
- TEX pipe can also read/write Register File
 - 3 SMs share 1 TEX
- Load/Store pipe can also read/write Register File




Constants

- Immediate address constants
- Indexed address constants
- Constants stored in memory, and cached on chip
 - L1 cache is per Streaming Multiprocessor




Shared Memory

- Each Stream Multiprocessor has 16KB of Shared Memory
 - 16 banks of 32bit words
- CUDA uses Shared Memory as shared storage visible to all threads in a thread block
 - Read and Write access



Execution Pipes

- Scalar MAD pipe
 - Float Multiply, Add, etc.
 - Integer ops,
 - Conversions
 - Only one instruction per clock
- Scalar SFU pipe
 - Special functions like Sin, Cos, Log, etc.
 - Only one operation per four clocks
- TEX pipe (external to SM, shared by all SM's in a TPC)
- Load/Store pipe
 - CUDA has both global and local memory access through Load/Store



GPGPU

What is really GPGPU?



- General Purpose computation using GPU in other applications than 3D graphics
 - GPU can accelerate parts of an application
- Parallel data algorithms using the GPUs properties
 - Large data arrays, streaming throughput
 - Fine-grain SIMD parallelism
 - Fast floating point (FP) operations
- Applications for GPGPU
 - Game effects (physics) nVIDIA PhysX
 - Image processing (Photoshop CS4)
 - Video Encoding/Transcoding (Elemental RapidHD)
 - Distributed processing (Stanford Folding@Home)
 - RAID6, AES, MatLab, etc.


Performance?

- Let's look at Standfords Folding@Home....
- Distributed Computing
- Folding@Home client is available for CUDA
 - Windows
 - All CUDA-enabled GPUs
- Performance GFLOPS:
 - Cell: 28
 - nVIDIA GPU: 110
 - ATI GPU: 109

| OS Type | Current TFLOPS* | Active CPUs | Total CPUs |
|------------------|-----------------|-------------|------------|
| Windows | 210 | 220936 | 2205693 |
| Mac OS X/PowerPC | 6 | 7316 | 120911 |
| Mac OS X/Intel | 24 | 7860 | 65496 |
| Linux | 35 | 32245 | 335464 |
| ATI GPU | 507 | 4612 | 12284 |
| nVIDIA GPU | 1614 | 14675 | 35626 |
| PLAYSTATION3 | 1679 | 59325 | 631489 |
| Total | 4095 | 347169 | 3406983 |

Previous GPGPU use, and limitations

- Working with a Graphics API
 - Special cases with an API like Microsoft Direct3D or OpenGL
- Addressing modes
 - Limited by texture size
- Shader capabilities
 - Limited outputs of the available shader programs
- Instruction sets
 - No integer or bit operations
- Communication is limited
 - Between pixels

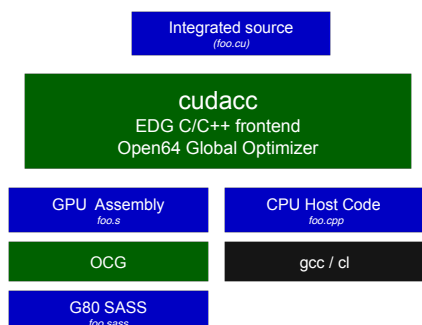


nVIDIA CUDA



- "Compute Unified Device Architecture"
- General purpose programming model
 - User starts several batches of threads on a GPU
 - GPU is in this case a dedicated super-threaded, massively data parallel co-processor
- Software Stack
 - Graphics driver, language compilers (Toolkit), and tools (SDK)
- Graphics driver loads programs into GPU
 - All drivers from nVIDIA now support CUDA.
 - Interface is designed for computing (no graphics ☺)
 - "Guaranteed" maximum download & readback speeds
 - Explicit GPU memory management

"Extended" C



Outline

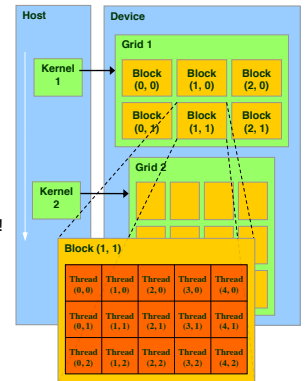
- The CUDA Programming Model
 - Basic concepts and data types
- The CUDA Application Programming Interface
 - Basic functionality
- More advanced CUDA Programming
 - 24th of October

The CUDA Programming Model

- The GPU is viewed as a compute **device** that:
 - Is a coprocessor to the CPU, referred to as the **host**
 - Has its own DRAM called **device memory**
 - Runs **many threads in parallel**
- Data-parallel parts of an application are executed on the device as **kernels**, which run in parallel on many threads
- Differences between GPU and CPU threads
 - GPU threads are extremely lightweight
 - Very little creation overhead
 - GPU needs 1000s of threads for full efficiency
 - Multi-core CPU needs only a few


Thread Batching: Grids and Blocks

- A kernel is executed as a **grid of thread blocks**
 - All threads share data memory space
- A **thread block** is a batch of threads that can **cooperate** with each other by:
 - Synchronizing their execution
 - Non synchronous execution is very bad for performance!
 - Efficiently sharing data through a low latency **shared memory**
- Two threads from two different blocks cannot cooperate



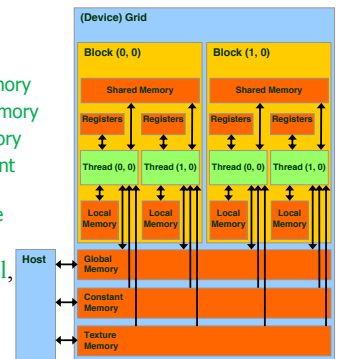
Block and Thread IDs

- Threads and blocks have IDs
 - Each thread can decide what data to work on
 - Block ID: 1D or 2D
 - Thread ID: 1D, 2D, or 3D
- Simplifies memory addressing when processing multidimensional data
 - Image and video processing (e.g. MJPEG...)



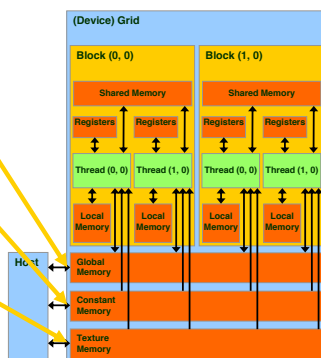
CUDA Device Memory Space Overview

- Each thread can:
 - R/W per-thread **registers**
 - R/W per-thread **local memory**
 - R/W per-block **shared memory**
 - R/W per-grid **global memory**
 - Read only per-grid **constant memory**
 - Read only per-grid **texture memory**
- The host can R/W **global, constant, and texture** memories



Global, Constant, and Texture Memories

- Global memory:
 - Main means of communicating R/W Data between **host** and **device**
 - Contents visible to all threads
- Texture and Constant Memories:
 - Constants initialized by host
 - Contents visible to all threads



Terminology Recap

- device = GPU = Set of multiprocessors
- Multiprocessor = Set of processors & shared memory
- Kernel = Program running on the GPU
- Grid = Array of thread blocks that execute a kernel
- Thread block = Group of SIMD threads that execute a kernel and can communicate via shared memory

| Memory | Location | Cached | Access | Who |
|----------|----------|----------------|------------|------------------------|
| Local | Off-chip | No | Read/write | One thread |
| Shared | On-chip | N/A - resident | Read/write | All threads in a block |
| Global | Off-chip | No | Read/write | All threads + host |
| Constant | Off-chip | Yes | Read | All threads + host |
| Texture | Off-chip | Yes | Read | All threads + host |

Access Times

- Register – Dedicated HW – Single cycle
- Shared Memory – Dedicated HW – Single cycle
- Local Memory – DRAM, no cache – “Slow”
- Global Memory – DRAM, no cache – “Slow”
- Constant Memory – DRAM, cached, 1...10s...100s of cycles, depending on cache locality
- Texture Memory – DRAM, cached, 1...10s...100s of cycles, depending on cache locality

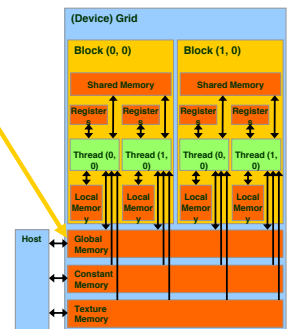
CUDA Highlights

- The API is an **extension to the ANSI C programming language**
 - Low learning curve than OpenGL/Direct3D
- The hardware is designed to enable lightweight runtime and driver
 - High performance

CUDA – API

CUDA Device Memory Allocation

- cudaMalloc()**
 - Allocates object in the device **Global Memory**
 - Requires two parameters
 - Address of a pointer** to the allocated object
 - Size of** allocated object
- cudaFree()**
 - Frees object from device Global Memory
 - Pointer to the object



CUDA Device Memory Allocation

- Code example:
 - Allocate a 64 * 64 single precision float array
 - Attach the allocated storage to Md.elements
 - “d” is often used to indicate a device data structure

```

BLOCK_SIZE = 64;
Matrix Md
int size = BLOCK_SIZE * BLOCK_SIZE * sizeof(float);

```

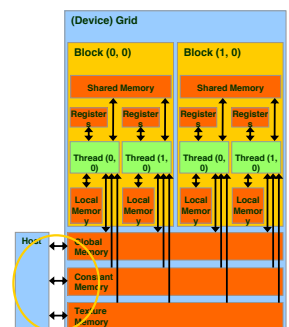
```

cudaMalloc((void**)&Md.elements, size);
cudaFree(Md.elements);

```

CUDA Host-Device Data Transfer

- cudaMemcpy()**
 - memory data transfer
 - Requires four parameters
 - Pointer to source
 - Pointer to destination
 - Number of bytes copied
 - Type of transfer
 - Host to Host
 - Host to Device
 - Device to Host
 - Device to Device
- Asynchronous in CUDA 1.3



Memory Management

- Device memory **allocation**
 - `cudaMalloc()`, `cudaFree()`
- Memory **copy** from host to device, device to host, device to device
 - `cudaMemcpy()`, `cudaMemcpy2D()`,
`cudaMemcpyToSymbol()`, `cudaMemcpyFromSymbol()`
- Memory **addressing**
 - `cudaGetSymbolAddress()`

CUDA Host-Device Data Transfer

- Code example:
 - Transfer a $64 * 64$ single precision float array
 - M is in host memory and Md is in device memory
 - `cudaMemcpyHostToDevice` and `cudaMemcpyDeviceToHost` are symbolic constants

```
cudaMemcpy(Md.elements, M.elements, size,  
           cudaMemcpyHostToDevice);
```

```
cudaMemcpy(M.elements, Md.elements, size,  
           cudaMemcpyDeviceToHost);
```

CUDA Function Declarations

| | Executed on the: | Only callable from the: |
|--|------------------|-------------------------|
| <code>__device__</code> float DeviceFunc() | device | device |
| <code>__global__</code> void KernelFunc() | device | host |
| <code>__host__</code> float HostFunc() | host | host |

- `__global__` defines a kernel function
 - Must return `void`
- `__device__` and `__host__` can be used together

CUDA Function Declarations

- `__device__` functions cannot have their address taken
- Limitations for functions executed on the device:
 - No recursion
 - No static variable declarations inside the function
 - No variable number of arguments

Calling a Kernel Function

- A kernel function must be called with an execution configuration:

```
__global__ void KernelFunc(...);  
dim3 DimGrid(100, 50); // 5000 thread blocks  
dim3 DimBlock(4, 8, 8); // 256 threads per block  
size_t SharedMemBytes = 64; // 64 bytes of shared memory  
KernelFunc <<< DimGrid, DimBlock, SharedMemBytes >>> (...);
```

- Any call to a kernel function is asynchronous from CUDA 1.0 on, explicit synch needed for blocking

Some Information on Toolkit

Compilation

- Any source file containing CUDA language extensions must be compiled with **nvcc**
- nvcc is a **compiler driver**
 - Works by invoking all the necessary tools and compilers like cudacc, g++, etc.
- nvcc can output:
 - Either C code
 - That must then be compiled with the rest of the application using another tool
 - Or object code directly

Linking

- Any executable with CUDA code requires two dynamic libraries:
 - The CUDA runtime library (**cudart**)
 - The CUDA core library (**cuda**)

Debugging Using Device Emulation

- An executable compiled in **device emulation mode** (**nvcc -deviceemu**) runs completely on the host using the CUDA runtime
 - No need of any device and CUDA driver
 - Each device thread is emulated with a host thread
- When running in device emulation mode, one can:
 - Use host native debug support (breakpoints, inspection, etc.)
 - Access any device-specific data from host code and vice-versa
 - Call any host function from device code (e.g. `printf`) and vice-versa
 - Detect deadlock situations caused by improper usage of `__syncthreads`

Lab Setup

- frogner.ndlab.net
 - GeForce 8600GT 256MB (G84)
 - 4 Multiprocessors, 32 Cores
- majorstuen.ndlab.net
 - GeForce 8600GT 512MB (G84)
 - 4 Multiprocessors, 32 Cores
- uranienborg.ndlab.net
 - GeForce 8600GT 512MB (G84)
 - 4 Multiprocessors, 32 Cores
- montebello.ndlab.net
 - GeForce 8800GT 256MB (G92)
 - 14 Multiprocessors, 112 Cores

Before you start...

- Four lines have to be added to your group users `.bash_profile` file

```
PATH=$PATH:/usr/local/cuda/bin
LD_LIBRARY_PATH=$LD_LIBRARY_PATH:/usr/local/cuda/lib

export PATH
export LD_LIBRARY_PATH
```
- When you use a machine, remember to update the message of the day! (etc/motd)

Compile and test SDK

- SDK is downloaded in the **/opt/** folder
- Copy and build in your users home directory
- Test machines uses Fedora Core 9, with gcc 4.3, SDK is for Fedora Core 8, and some fixing is needed to compile...

Add the following **#include** in these files:

```
common/src/paramgl.cpp: <cstring>
projects/cppIntegration/main.cpp: <cstdlib>
common/inc/exception.h: <cstdlib>
common/inc/cutil.h: <cstring>
common/inc/cmd_arg_reader.h: <typeinfo>
```


Some usefull resources

nVIDIA CUDA Programming Guide 2.0

http://developer.download.nvidia.com/compute/cuda/2_0/docs/NVIDIA_CUDA_Programming_Guide_2.0.pdf

nVIDIA CUDA Reference Manual 2.0

http://developer.download.nvidia.com/compute/cuda/2_0/docs/CudaReferenceManual_2.0.pdf

nVISION08: Getting Started with CUDA

http://www.nvidia.com/content/cudazone/download/Getting_Started_w_CUDA_Training_NVISION08.pdf